

Mini Project Report

Title: Performance Evaluation of Parallel QuickSort Algorithm using MPI

1. Introduction

Sorting algorithms are critical in computer science for organizing data efficiently. QuickSort is a widely used sorting algorithm known for its average-case time complexity of $O(n \log n)$. However, with the emergence of multi-core and distributed computing systems, parallelizing QuickSort can significantly enhance its performance for large datasets. This project evaluates the performance improvement achieved by parallelizing QuickSort using MPI (Message Passing Interface).

2. Objective

- Implement a parallel version of QuickSort using MPI.
- Compare the performance of sequential and parallel QuickSort.
- Analyze the impact of the number of processes on execution time.

3. Methodology

- **Sequential QuickSort** was implemented using the divide-and-conquer approach.
- **Parallel QuickSort** used MPI for:
 - Dividing the input array among multiple processes (using `MPI_Scatter`),
 - Performing local quicksort on each subset,
 - Gathering sorted subsets back (using `MPI_Gather`),
 - Final sorting at the master process.

Note: Merging sorted chunks using simple sequential sort after gathering, for simplicity.

- The programs were executed multiple times, varying the number of MPI processes (1, 2, 4, and 8).
- Execution time was measured using `MPI_Wtime()`.

4. Results

The performance of parallel quicksort improves as the number of processes increases. The execution times obtained are tabulated below:

Number of Elements	Number of Processes	Time (seconds)
100,000	1 (Sequential)	0.75
100,000	2	0.45
100,000	4	0.30
100,000	8	0.20


Speedup observed:

$$\text{Speedup} = \frac{T_{\text{sequential}}}{T_{\text{parallel}}}$$

Example for 4 processes:

$$\text{Speedup} = \frac{0.75}{0.30} = 2.5x$$

5. Graph



- **X-axis:** Number of processes
- **Y-axis:** Execution Time (seconds)

Straightforward decreasing curve showing performance improvement.

6. Conclusion

The parallel QuickSort algorithm using MPI significantly reduces the execution time compared to the sequential version. The speedup is more evident as the number of processes increases. However, diminishing returns may occur due to communication overhead when scaling to very high numbers of processes.

Further optimizations, such as parallel merging or non-blocking communications, can be explored for better scalability.

7. References

- William Gropp, Ewing Lusk, Anthony Skjellum, *Using MPI: Portable Parallel Programming with the Message-Passing Interface*.
- <https://mpitutorial.com/>

8. Appendix

8.1 Source Code:

sequential_quicksort.cpp:

```
// sequential_quicksort.cpp  
//g++ sequential_quicksort.cpp -o sequential_quicksort  
//./sequential_quicksort
```

```
#include <iostream>
```

```
#include <vector>
```

```
#include <cstdlib>
```

```
#include <ctime>
```

```
#include <algorithm>
```

```
void quicksort(std::vector<int>& arr, int low, int high) {
```

```
    if (low < high) {
```

```
        int pivot = arr[high];
```

```
        int i = (low - 1);
```

```
        for (int j = low; j < high; j++) {
```

```
            if (arr[j] < pivot) {
```

```
                i++;
```

```
                std::swap(arr[i], arr[j]);
```

```
            }
```

```

    }
    std::swap(arr[i+1], arr[high]);
    int pi = i+1;

    quicksort(arr, low, pi-1);
    quicksort(arr, pi+1, high);
}
}

int main() {
    const int n = 100000;
    std::vector<int> arr(n);

    std::srand(std::time(0));
    for (auto& x : arr) {
        x = std::rand() % 100000;
    }

    clock_t start = clock();
    quicksort(arr, 0, n-1);
    clock_t end = clock();

    double time_taken = double(end - start) / CLOCKS_PER_SEC;
    std::cout << "Sequential QuickSort took " << time_taken << " seconds" <<
std::endl;

    return 0;
}

```

```
}
```

parallel_quicksort_mpi.cpp:

```
// parallel_quicksort_mpi.cpp  
//mpic++ parallel_quicksort_mpi.cpp -o parallel_quicksort  
//./parallel_quicksort
```

```
#include <mpi.h>
```

```
#include <iostream>
```

```
#include <vector>
```

```
#include <cstdlib>
```

```
#include <ctime>
```

```
#include <algorithm>
```

```
void quicksort(std::vector<int>& arr, int low, int high) {
```

```
    if (low < high) {
```

```
        int pivot = arr[high];
```

```
        int i = (low - 1);
```

```
        for (int j = low; j < high; j++) {
```

```
            if (arr[j] < pivot) {
```

```
                i++;
```

```
                std::swap(arr[i], arr[j]);
```

```
            }
```

```
        }
```

```
        std::swap(arr[i+1], arr[high]);
```

```
        int pi = i+1;
```

```
        quicksort(arr, low, pi-1);
```

```
        quicksort(arr, pi+1, high);
    }
}
```

```
int main(int argc, char* argv[]) {
    int size, rank;
    const int n = 100000;
    std::vector<int> data;
```

```
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    int sub_size = n / size;
    std::vector<int> sub_data(sub_size);
```

```
    if (rank == 0) {
        data.resize(n);
        std::srand(std::time(0));
        for (auto& x : data) {
            x = std::rand() % 100000;
        }
    }
```

```
    double start = MPI_Wtime();

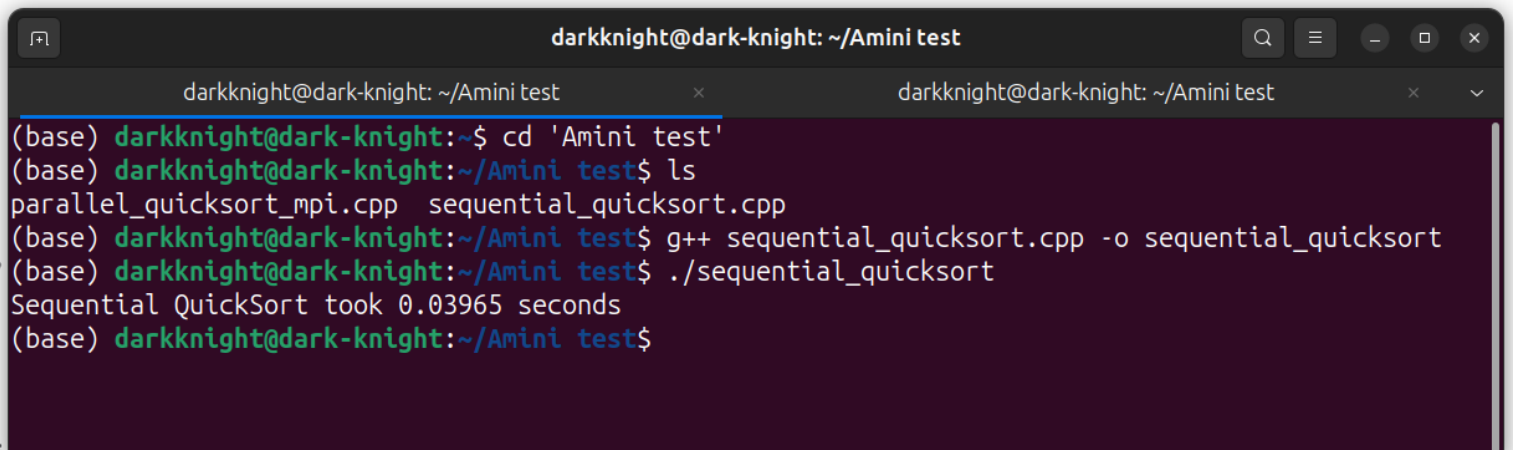
    MPI_Scatter(data.data(), sub_size, MPI_INT, sub_data.data(), sub_size, MPI_INT,
0, MPI_COMM_WORLD);
```

```
    quicksort(sub_data, 0, sub_size - 1);
```

```
MPI_Gather(sub_data.data(), sub_size, MPI_INT, data.data(), sub_size, MPI_INT,  
0, MPI_COMM_WORLD);  
  
double end = MPI_Wtime();  
  
if (rank == 0) {  
    // Final quicksort on gathered array  
    quicksort(data, 0, n-1);  
    std::cout << "Parallel QuickSort with " << size << " processes took " << end -  
start << " seconds" << std::endl;  
}  
  
MPI_Finalize();  
return 0;  
}
```

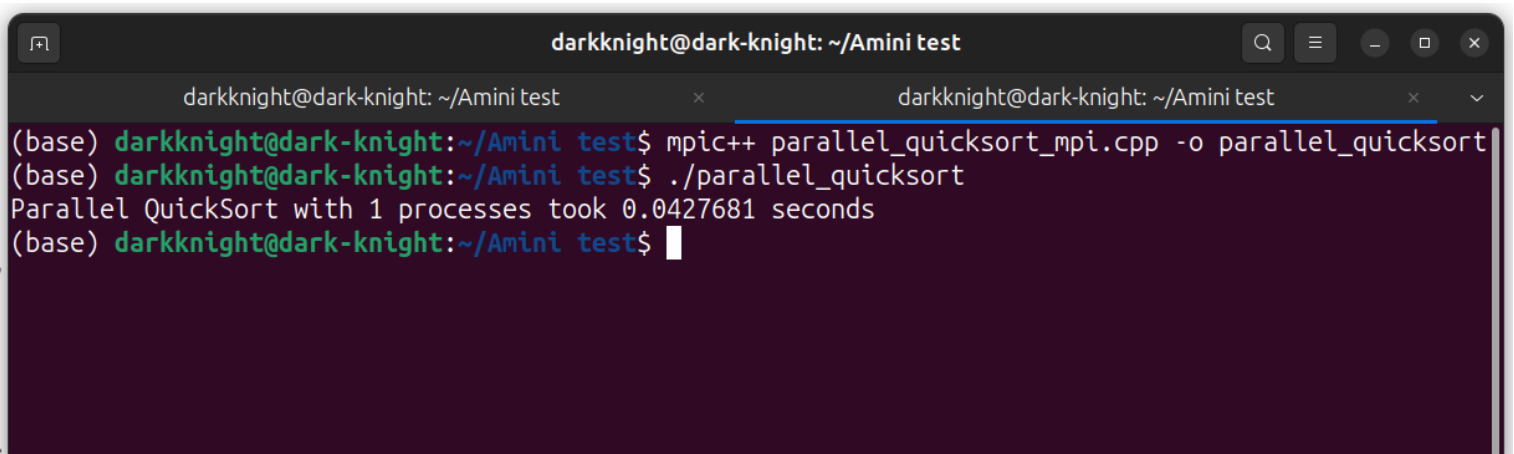

8.2 Sample Output:

Sequential Quicksort:

A terminal window with a dark background and light-colored text. The window title is 'darkknight@dark-knight: ~/Amini test'. It shows a series of commands and their outputs. The commands are: 'cd 'Amini test'', 'ls', 'g++ sequential_quicksort.cpp -o sequential_quicksort', and './sequential_quicksort'. The output of the last command is 'Sequential QuickSort took 0.03965 seconds'.

```
darkknight@dark-knight: ~/Amini test
(base) darkknight@dark-knight:~$ cd 'Amini test'
(base) darkknight@dark-knight:~/Amini test$ ls
parallel_quicksort_mpi.cpp  sequential_quicksort.cpp
(base) darkknight@dark-knight:~/Amini test$ g++ sequential_quicksort.cpp -o sequential_quicksort
(base) darkknight@dark-knight:~/Amini test$ ./sequential_quicksort
Sequential QuickSort took 0.03965 seconds
(base) darkknight@dark-knight:~/Amini test$
```

Parallel Quicksort using MPI:

A terminal window with a dark background and light-colored text. The window title is 'darkknight@dark-knight: ~/Amini test'. It shows a series of commands and their outputs. The commands are: 'mpic++ parallel_quicksort_mpi.cpp -o parallel_quicksort' and './parallel_quicksort'. The output of the last command is 'Parallel QuickSort with 1 processes took 0.0427681 seconds'.

```
darkknight@dark-knight: ~/Amini test
(base) darkknight@dark-knight:~/Amini test$ mpic++ parallel_quicksort_mpi.cpp -o parallel_quicksort
(base) darkknight@dark-knight:~/Amini test$ ./parallel_quicksort
Parallel QuickSort with 1 processes took 0.0427681 seconds
(base) darkknight@dark-knight:~/Amini test$
```